

CS341: Computer Architecture Lab

Lab 5: Hacking ChampSim

Report

Team ARCHimedes

Vibhav Aggarwal (190050128)

Adithya Bhaskar (190050005)

Devansh Jain (190100044)

Harshit Varma (190100055)

Harshit Gupta (190050048)



Department of Computer Science and Engineering
Indian Institute of Technology Bombay

2021-2022

Contents

1	Understanding the problem statement	2
1.1	Summary of MadCache	2
1.2	Understanding ChampSim code	3
2	LLC Bypassing	5
2.1	Implementation	5
2.2	Source files edited or added	8
2.3	Impact on IPC	8
3	Prefetching	9
3.1	Implementation	9
3.1.1	Throttle Policy	9
3.1.2	Selecting d_l, d_u	9
3.2	Normalized IPC improvement wrt thresholds	11
3.3	Prefetcher accuracies and coverage wrt thresholds	12
3.4	Results (Values)	13
3.5	Source files edited or added	14
4	Running the code	15
4.1	LLC Bypassing	15
4.2	Prefetching	15
5	Contributions	16

Abstract

This lab consists of two parts, LLC Bypassing and Prefetcher Filtering.

In the first part of the lab, we implement an Instruction Pointer/Program Counter-based adaptive insertion policy for LLC in ChampSim. We refer to a paper on MadCache, an example of such a policy, and implement a similar policy. Our implementation increases the IPC by roughly 20% for the given trace file.

In the second part of the lab, we implement prefetch filtering/throttling for a simple IP-based prefetcher. Our implementation increases the IPC by roughly 16% for the given trace file.

1. Understanding the problem statement

1.1 Summary of MadCache

MadCache is an adaptive insertion policy that uses the Program Counter (PC) behavior history to determine the appropriate insertion policy for the Last Level Cache (LLC). It decides whether a miss should be saved in the cache or bypassed to the L2 Cache (L2C). In other words, choose between LRU (Least Recently Used) mode and bypassing mode. The underlying assumption is that if a particular PC in a program exhibits streaming behavior, it will likely continue to do so, and should be prevented from flushing out other useful entries in the LLC. While a PC may exhibit streaming behavior for one portion of the program, it is possible it changes its behavior later on. For such a case, thrashing protection must be enabled to allow a PC to get out of bypass mode. Bypassing Bimodal Insertion Policy is the Bypassing version of BIP, which bypasses the majority of the incoming cachelines anticipating streaming activity and inserts a few lines to the LLC.

MadCache uses four structures to adapt the insertion policy of the LLC. A 10-bit counter whose MSB is used to determine the default insertion policy. If the MSB is 1, then the default policy is BBIP. A subset of the cache sets are called trackers sets, and the behavior of PCs accessing this set determines the counter value through set dueling. The rest of the sets, in fact, the majority of the cache are called follower sets, and they simply adhere to the default insertion policy. A lookup table referred to as the “PC-predictor” is used to keep track of the history of PCs that have accessed the tracker sets. It has 1024 entries of 3 columns - (Default policy at the time of the entry of this PC, PC), a 6-bit counter whose MSB determines the behavior for this PC, and the number of entries associated with the PC.

At the start, the PC-predictor and the cache are empty, and the counter is set to 511, i.e., just below the threshold for BBIP. After an initial LLC miss for a cacheline in the tracker sets, the invoking PC would be added to the PC-predictor. The counter for the entry is set to 63, i.e., just below the threshold of BBIP. The cacheline would be added to the LLC. The cacheline (in the tracker set) would also store the index to the table and a reuse bit to reward or penalize the tracked PC entry. If the cacheline is accessed again, the counter for the PC entry would be decremented (in favor of LRU). However, if the cacheline is evicted without being reused, then the counter would be incremented (in favor of BBIP). Entries from the PC-predictor are evicted when the number of cachelines brought in by the PC entry falls to zero.

A more detailed explanation of the MadCache policy is given in chapter 2.

MadCache performs faster than both LRU and DIP, resulting in a 4.5% speedup over LRU.

1.2 Understanding ChampSim code

- The `handle_read` function is responsible for reading a query from the read queue and processing it.

A read hit means that the line requested is present in the cache. So, the data is read from the cache itself. Following are the things that are done on a hit:

- Get a read request from the read queue `RQ`. If no read request, return.
 - Get the set and the way corresponding to the request address. For a hit, there must be a way with same address as the requested one. This is done using the `get_set` and `check_hit` functions.
 - Read the data and update it in the packet that we got from the read queue.
 - If it is a load request, update the prefetcher using the `prefetcher_operate` functions
 - Update the replacement policy using the `update_replacement_policy` functions. This corresponds to bringing the accessed block to MRU position in the LRU scheme.
 - Mark the entry in the MSHR of higher level as completed and add the data to it. This is done using the `return_data` function for the correct upper level cache. Notice that for L1D, L1I, ITLB and DTLB do not have a higher level cache. Instead, they put the data in the PROCESSING queue.
 - Update the statistics like number of hits and number of accesses.
- A read miss means that the line requested is not present in the cache. In such a case, we do the following:(skipping over getting the query and detecting a miss as it is same as hit)
 - First we need to check if there is a read request for the same address already in MSHR or not. If there is we can merge the requests. This is done using the `check_mshr` function.
 - * If the address is not in MSHR, we need to add a new entry in MSHR and read queue of the lower level. Adding to MSHR is done using `add_mshr` function. Adding to the lower level's read queue is done using `lower_level->add_rq` function. (This function also checks if the request can be served from the write queue or not and adds the request to read queue only if it can't be done)
During the above operations, we need to ensure that we don't exceed the size of MSHR's and the queues. If this is the case, we must stall.
 - * If the address is in MSHR, we need to merge the requests.
There is a member set named `index_depend_on_me` (different sets for load, store and instruction) in the `PACKET` class. This set stores the index of all the requests depending on it. Merging indices amounts to joining the correct sets from request to MSHR entry and insert the index of request to the MSHR entry.
 - Update statistics like number of misses, number of MSHR entries merged etc.
- The `handle_fill` function is responsible for processing the completed requests in MSHR and put the data in correct position in the cache. Any function serving a read request uses `return_data` which only adds the data to the MSHR entry for that address. But we need the data in cache

and not in MSHR. The work of getting the data from MSHR and writing to cache is done by `handle_fill`. That is why it needs to read (completed) request from MSHR.

- The `lower_level`, `upper_level` data members are declared in the `MEMORY` class. Since `CACHE` class inherits from `MEMORY`, it also has these data members. They are assigned in `main.cc` in the `main` function for all the caches.

These are arrays which store pointer to the next lower or next higher cache in the memory hierarchy (if they do not exist, then set to `NULL`). Since we have different instruction and data cache at L1, we need 2 different `upper_level` cache array, one for data and one for instruction. So, for example for L2 cache, the values are:

```
ooo_cpu[i].upper_level_icache = ooo_cpu[i].L1I
ooo_cpu[i].upper_level_dcache = ooo_cpu[i].L1D
ooo_cpu[i].lower_level = uncore.LLC
```

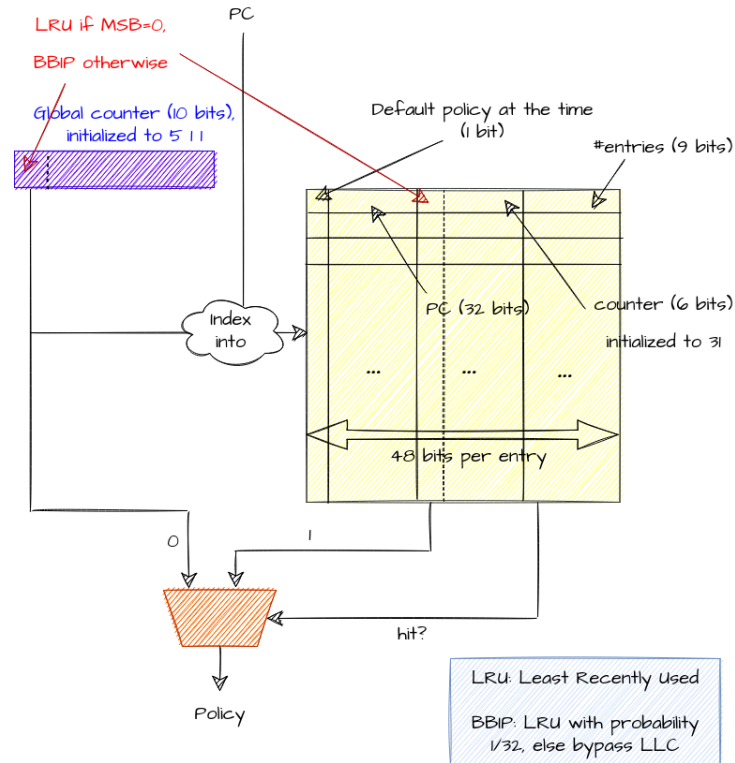
Notice that L1D, L1I, L2 and the TLBs are CPU specific while the LLC and DRAM are shared.

- The `return_data` function is responsible for adding the data in an MSHR entry and marking it as completed. It takes a `PACKET` object as argument which stores the address and data. The basic workflow is:
 - Find the MSHR entry with the same address as the one mentioned in the given packet. This is done using `check_mshr` function. If it does not exist, return error
 - Mark the MSHR entry as completed (setting `returned` to `COMPLETED`) and add the data and metadata to the MSHR entry.
 - Update the metrics like number of requests returned and latency
 - Update `next_fill_index` using `update_fill_cycle` function. This index is the index from which `read_fill` reads the request to be written in the cache.

2. LLC Bypassing

2.1 Implementation

The relevant output files are in their respective directories, as required by the problem statement. We proceed to describe our implementation of MadCache. We begin with the description of the *predictor* structure. Refer to the illustration below.

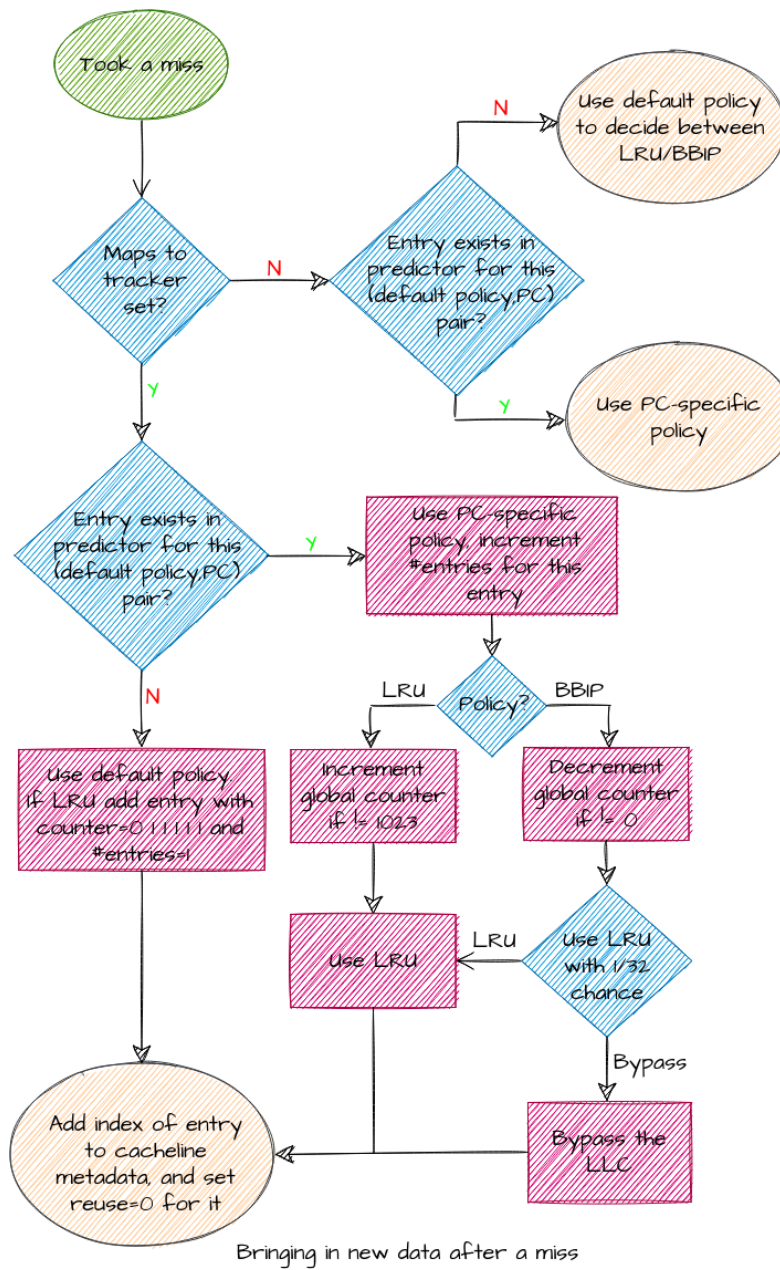


The structure of the predictor table

The predictor consists of a global 10-bit counter (which determines the global default policy, initialized to 0111111111 or 511) and a set of entries each with their counters, which determines PC-specific policies. In particular, if the MSB of the global counter is 0, the default policy is LRU; else, it is to use BBIP. Here, BBIP uses LRU 1/32-th of the time and bypasses otherwise. Each entry contains:

- **Default policy and PC:** This consists of the PC used when fetching the cacheline which led to the creation of this entry, and the default policy at the time.
- **Counter:** A 6-bit counter governing the PC-specific policy, of 6 bits.
- **Number of entries:** This field tells us how many cachelines currently use this entry (0 if none).

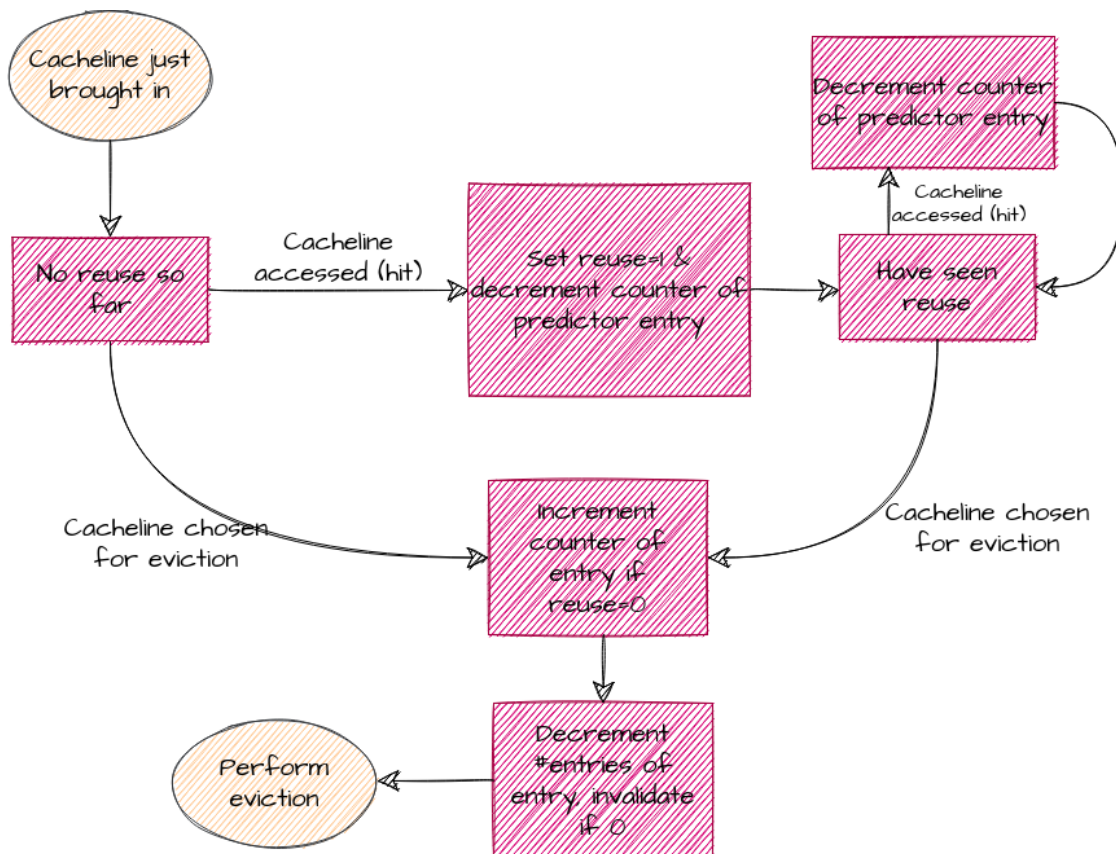
Next we describe how we use this predictor to decide what to do with newly fetched cachelines from DRAM to the LLC. Refer to the flowchart below.



If the (Policy, PC) pair exists in the predictor, we use the custom policy to decide what to do. This occurs even for lines not in tracker sets (we map those cachelines whose index bits have 0101 as the 4 LSBs as tracker sets; i.e., 1/16 of all cachelines are tracker sets). However, for these non-tracker lines, we do **not** create an entry in the predictor. However, if the cacheline *is* part of a tracker set, we check to see if an entry exists in the predictor. If so, we use the corresponding custom policy and increment its #entries field if it happens to be LRU (we shouldn't and don't do this for Bypass since there is no cacheline allocated in that case). Otherwise, we fall back to using the default policy and try to allocate a free predictor entry for this cacheline. If this succeeds, we initialize it with a counter of 011111 ("just" LRU) and a #entries field of 1. Note that if the allocation fails, we treat this cacheline just as if it were a non-tracker line henceforth (until new data is brought into it again).

If the policy used is LRU, we initialize the cacheline by setting reuse=0 and the pc_pred_index to point to the allocated predictor entry.

Further, whenever we see a miss that maps to a tracker set, we check if we have a PC-specific policy for this PC. If so, we correspondingly change the global counter to penalize that specific type of behavior (since it caused us a miss). This implements set dueling (or "PC dueling", rather) between the LRU and Bypass sets. We next look at the lifetime of a tracker cacheline in the LLC. Refer to the flowchart below.



Life-cycle of a cacheline in a tracker set

A tracker cacheline starts with a `reuse` bit of 0 and a counter initialized to 31 (if it was the first one for that entry). Whenever this cacheline sees a hit, it means that we were right to use LRU; thus, we decrement the corresponding predictor entry's counter to make the behavior more "LRU-like" for this PC. We also set `reuse=1` for this cacheline on a hit. On the other hand, when the cacheline is being evicted, if we see a reuse bit of 0, this means that the cacheline was never used again; we should have bypassed the LLC. Thus we increment the corresponding counter. In either case, we decrement the `#entries` field of the counter on eviction, and if it falls to 0, de-allocate the predictor entry (so that it can be used by future misses).

2.2 Source files edited or added

We modified the following files to implement the above scheme:

- The files `inc/block.h` and `src/block.cc` contain the implementation of the predictor class.
- The file `inc/cache.h` has the definitions of a few helper functions and the modifications of the `CACHE` class to include the predictor.
- The file `src/cache.cc` contains the bulk of the changes, and implements the core logic of the above scheme. In particular, the functions that handle reads, writebacks and fills see extensive changes. Some other functions also see minor modifications.

2.3 Impact on IPC

The baseline LRU scheme gave an IPC of 0.142776 on average. On the other hand, our scheme gave an IPC of 0.171424 on average.

This corresponds to an increase (improvement) in IPC by roughly **20.065%**.

3. Prefetching

3.1 Implementation

Each IP initially has the prefetch degree set to 3. We dynamically update it based on the per-IP accuracy of the prefetches.

For each IP, we track the number of useful and useless prefetches. Whenever any prefetch is marked as useful or useless in `cache.cc`, the `prefetch_throttle` function is called (depending on the cache type), and these values are updated for the respective IP. `prefetch_throttle` updates the degree based on the throttle policy.

3.1.1 Throttle Policy

For a cache, we define d_u, d_l as the upper and lower bounds on the prefetch degree and t as the threshold. For an IP, updates to the degree occur only after `prefetch_throttle` is called 10 times for that IP; this counter is reset after every update.

Updates happen based on the accuracy, which is the ratio of useful to useful + useless prefetches.

If accuracy is $> t$, then the degree is incremented by 1, unless it's already equal to d_u .

If accuracy is $\leq t$, then the degree is decremented by 1, unless it's already equal to d_l .

Sections 3.2 & 3.3 report the results for different values of t , for each cache, keeping d_l, d_u fixed as:

L1D: $d_l = 1, d_u = 8$

L2C: $d_l = 1, d_u = 16$

3.1.2 Selecting d_l, d_u

d_l has been fixed to 1 throughout; we don't tune this parameter. Figures 3.1, 3.2 contain the results that help us choose d_u . Section 3.4 contains the respective numeric values obtained.

For L1D, we first fix the parameters for L2C at the baseline ($d_u = d_l = 3$, i.e., no throttling) then we compute the metrics for all combinations of $t \in \{0.25, 0.5, 0.75\}$ and $d_u \in \{4, 6, 8\}$. We observe that $d_u = 8$ gives the highest IPC. Thus, we set $d_u = 8$ for L1D.

For L2C, we first fix the parameters for L1D at the baseline ($d_u = d_l = 3$, i.e., no throttling) then we compute the results for all combinations of $t \in \{0.25, 0.5, 0.75\}$ and $d_u \in \{6, 8, 10, 12\}$. From these results, we observe that the results are poor for threshold values 0.5 and 0.75 and stay approximately the same. This could be since the accuracy values for IPs in L2C are low; they always tend to stay below 0.5; thus, the degree quickly reduces to 1 and stays 1 throughout. This motivates us to change the set of t values that we are considering. We also observe that for a fixed threshold, the best results are obtained

for $d_u = 16$. Thus, we run more experiments for $d_u = 16$ and $t \in \{0.05, 0.15\}$. We observe that the best IPC is obtained for $t = 0.05$ and $d_u = 16$. Thus, we set $d_u = 16$ for L2C.

Finally, after fixing d_u, d_l for both the caches, we compute the results for all combinations of $t_{L1D} \in \{0.25, 0.5, 0.75\}$ and $t_{L2C} \in \{0.05, 0.15, 0.25\}$ and report these in sections 3.2 and 3.3.




Figure 3.1: Labels on the x -axis are of the form $t-d_u$




Figure 3.2: Labels on the x -axis are of the form $t-d_u$

3.2 Normalized IPC improvement wrt thresholds




Figure 3.3: Labels on the x -axis are of the form $t_{L1D}-t_{L2C}$

As can be seen from the plot, the best IPC improvement achieved is about **16%** for the thresholds 0.25 for L1D and 0.05 for L2C.

3.3 Prefetcher accuracies and coverage wrt thresholds




Figure 3.4: Labels on the x -axis are of the form $t_{L1D}-t_{L2C}$




Figure 3.5: Labels on the x -axis are of the form $t_{L1D}-t_{L2C}$

3.4 Results (Values)

-----L1D-----				
+-----+	+-----+	+-----+	+-----+	+-----+
t-d_u	Accuracy	Coverage	IPC	
-----+	-----+	-----+	-----+	-----+
0.25-4	0.712436	0.269298	0.222645	
0.25-6	0.600354	0.258107	0.243394	
0.25-8	0.652155	0.319944	0.247193	
0.5-4	0.712436	0.269298	0.222645	
0.5-6	0.691499	0.265962	0.23328	
0.5-8	0.661731	0.317649	0.24555	
0.75-4	0.786124	0.233487	0.21702	
0.75-6	0.754459	0.19353	0.211511	
0.75-8	0.754403	0.199547	0.212759	
+-----+	-----+	-----+	-----+	+-----+

Figure 3.6: Results related to figure 3.1. Result files included in `prefetch_throttling_result/L1D`

-----L2C-----				
+-----+	+-----+	+-----+	+-----+	+-----+
t-d_u	Accuracy	Coverage	IPC	
-----+	-----+	-----+	-----+	-----+
0.05-16	0.322317	0.364969	0.242624	
0.15-16	0.322932	0.365557	0.242304	
0.25-06	0.251878	0.221186	0.218375	
0.25-08	0.258495	0.229774	0.218577	
0.25-10	0.277876	0.256039	0.22223	
0.25-12	0.280375	0.260855	0.22239	
0.25-14	0.284907	0.267188	0.222789	
0.25-16	0.285274	0.267735	0.222762	
0.5-06	0.177547	0.13465	0.209062	
0.5-08	0.177547	0.13465	0.209062	
0.5-10	0.177679	0.134674	0.208561	
0.5-12	0.177547	0.13465	0.209062	
0.75-06	0.177679	0.134674	0.208561	
0.75-08	0.177679	0.134674	0.208561	
0.75-10	0.177679	0.134674	0.208561	
0.75-12	0.177679	0.134674	0.208561	
+-----+	-----+	-----+	-----+	+-----+

Figure 3.7: Results related to figure 3.2. Result files included in `prefetch_throttling_result/L2C`

-----L1D-----				
t_1d-t_2c	Accuracy	Coverage	IPC	
0.25-0.05	0.653101	0.331389	0.247562	
0.25-0.15	0.655644	0.322231	0.23995	
0.25-0.25	0.652448	0.308576	0.23507	
0.5-0.05	0.676553	0.321961	0.247528	
0.5-0.15	0.671388	0.315496	0.239972	
0.5-0.25	0.666546	0.305828	0.233228	
0.75-0.05	0.77162	0.196809	0.238049	
0.75-0.15	0.77177	0.196702	0.237963	
0.75-0.25	0.77263	0.196597	0.231663	

Figure 3.8: Results related to figure 3.4. Files included in `prefetch_throttling_result/combined`

-----L2C-----				
t_1d-t_2c	Accuracy	Coverage	IPC	
0.25-0.05	0.241089	0.319648	0.247562	
0.25-0.15	0.220957	0.27399	0.23995	
0.25-0.25	0.190415	0.215303	0.23507	
0.5-0.05	0.251362	0.324468	0.247528	
0.5-0.15	0.228838	0.27902	0.239972	
0.5-0.25	0.192552	0.208844	0.233228	
0.75-0.05	0.371984	0.35425	0.238049	
0.75-0.15	0.371986	0.354375	0.237963	
0.75-0.25	0.358576	0.32765	0.231663	

Figure 3.9: Results related to figure 3.5. Files included in `prefetch_throttling_result/combined`

3.5 Source files edited or added

Edited:

ChampSim/inc/cache.h
ChampSim/src/cache.cc
ChampSim/prefetcher/ip_stride.l2c_pref
ChampSim/prefetcher/no.l2c_pref
ChampSim/prefetcher/no.l1d_pref

Added:

ChampSim/inc/ip_stride.h
ChampSim/prefetcher/ip_stride.l1d_pref

4. Running the code

4.1 LLC Bypassing

The decision variable used to decide whether to use the implemented version or the baseline is `LLC_BYPASS` in `champsim.h` file. If it is defined, then the code will use Madcache; otherwise, it uses the default LRU policy.

The commands used are:

```
./build_champsim.sh bimodal no no no no lru 1
./run_champsim.sh bimodal-no-no-no-no-lru-1core 10 10 trace.champsimtrace.xz
```

4.2 Prefetching

Disable LLC Bypassing as explained above.(Comment out the line defining `LLC_BYPASS` in `champsim.h`.) The prefetchers are defined in the files `ip_stride.l2c_pref` and `ip_stride.l1d_pref`. So, depending on which prefetcher to use, pass the name while running `build_champsim.sh`.

The commands will be:

Only L1D prefetching:

```
./build_champsim.sh bimodal no ip_stride no no lru 1
./run_champsim.sh bimodal-no-ip_stride-no-no-lru-1core 10 10 trace.champsimtrace.xz
```

Only L2C prefetching:

```
./build_champsim.sh bimodal no no ip_stride no lru 1
./run_champsim.sh bimodal-no-no-ip_stride-no-lru-1core 10 10 trace.champsimtrace.xz
```

Combined:

```
./build_champsim.sh bimodal no ip_stride ip_stride no lru 1
./run_champsim.sh bimodal-no-ip_stride-ip_stride-no-lru-1core 10 10 trace.champsimtrace.xz
```

The threshold and the bounds are defined in the respective file with the names `DEGREE_LOWER_BOUND`, `DEGREE_UPPER_BOUND` and `ACCURACY_TH`.

Set both `DEGREE_LOWER_BOUND` and `DEGREE_UPPER_BOUND` to 3 to get the baseline result.

5. Contributions

Table 5.1: Contributions of each team member

Member	Work Done
Vibhav Aggarwal	Part 3 code and experiments, plotter scripts
Adithya Bhaskar	Part 1 code, part 2 explanation
Devansh Jain	Part 0 summary, part 1 code testing, result for part 2
Harshit Varma	Part 3 code and experiments, part 4
Harshit Gupta	Part 0 code explanation, part 1 code testing, code merging